Population-scale coupling between electrotonic regime and synaptic spatial organization in mouse visual cortex

Pretham Voruganti
Independent Researcher

Abstract

Dendritic branches differ in their capacity to electrically isolate synaptic inputs, ranging from electrotonically compact domains where inputs sum linearly to compartmentalized domains capable of local nonlinear integration. Whether this biophysical heterogeneity is reflected in the spatial arrangement of synapses remains poorly understood, in part because prior evidence has been limited to small samples and shallow mechanistic tests. Here we report a population-scale analysis of 2,138 proofread neurons (136,414 dendritic branches, ∼8.2 million synapses) from the MICrONS mouse visual cortex electron microscopy volume. Using continuous electrotonic length (L/λ) as a predictor in mixed-effects models, we find that synaptic spatial organization covaries with dendritic biophysical regime: electrotonically isolated branches carry more structured synaptic arrangements, as measured by two independent spacing metrics (Clark-Evans z: inhibitory $\beta = 7.65$, $p < 10^{-77}$; Interval CV z: inhibitory $\beta = -17.5$, $p < 10^{-90}$). This coupling is robust to threshold choice (8/8 discretization schemes converge) and is not explained by excitatory/inhibitory input balance (mediation $< 7\%$). Partner architecture features — partner diversity and contact multiplicity — mediate 25–53% of the coupling at population scale, supporting two separable organizational channels: partner diversity tracks local spacing regularity while contact multiplicity tracks global synapse dispersion. Random-slope mixed models reveal that synaptic spacing organization is structured by cell subtype rather than a binary excitatory/inhibitory divide. Excitatory neurons show steep $L/\lambda \rightarrow$ spacing regularity slopes (population median $= -2.23$), with L6-CT corticothalamic neurons showing the strongest coupling (median $= -3.04$). Inhibitory subtypes are bimodal: BC (PV-like) neurons show an inverted profile (median $= +1.81$), consistent with their proximal-domain specialization, while BPC and NGC neurons show slopes steeper than the excitatory average (-2.76 and -3.25). The pooled inhibitory median (-0.07) is an averaging artifact masking this heterogeneity (Kruskal-Wallis $p = 2.7 \times 10^{-11}$ across subtypes). We decompose the inhibitory residual — the ∼86% of inhibitory coupling not explained by partner architecture — by tracing afferent cell type identity through a multi-stage pathway. Basket cell (BC) afferents are strongly enriched on compact branches (neuron-level median $\log_2$ enrichment $= 0.99$, 89% of neurons positive, $p = 4.0 \times 10^{-45}$), concentrated in proximal dendritic compartments. Afferent identity and partner architecture jointly explain 18–26% of the inhibitory regime effect, reducing the residual from ∼86% to ∼75%. Critically, within-axon bouton placement analysis reveals that BC contacts are placed more compactly on compact branches even after conditioning on contact multiplicity (regime coefficient $= +0.95$, $p = 1.1 \times 10^{-107}$; equal-k subsets: $p < 10^{-18}$ at every k), and that Martinotti cell (MC) contacts spread more broadly than BC contacts on compartmentalized branches (interaction $p = 5.2 \times 10^{-4}$). These results trace a multi-stage pathway — electrotonic regime $\rightarrow$ afferent subtype composition $\rightarrow$ partner architecture $\rightarrow$ within-axon placement $\rightarrow$ synaptic spacing — that partially decomposes the inhibitory residual and supports subtype-specific placement rules as a previously uncharacterized layer of inhibitory circuit organization.

1 Introduction

Dendritic trees are not passive conduits. The electrotonic properties of a branch — its length, diameter, and distance from the soma — determine how synaptic inputs are integrated, from linear summation in compact domains to supralinear amplification in electrically isolated compartments (Rall, 1967; London and Häusser, 2005; Koch, 1999). This biophysical heterogeneity has long been recognized as the basis for dendritic computation: thin, isolated branches can function as independent computational subunits, generating local NMDA spikes that are thresholded before reaching the soma (Mel, 1994; Polsky et al., 2004; Poirazi et al., 2003; Larkum et al., 2009).

Most connectomics studies characterize synaptic connectivity in terms of *who* connects to *whom* — connection probabilities, motif statistics, cell-type circuit diagrams. A complementary question, less frequently addressed, is what generative rules produce the *spatial* architecture of connections on dendrites. Spatial statistics on synapse positions can, in principle, reveal these rules: if the arrangement of synapses varies systematically with the biophysical properties of the branch, and if that variation can be decomposed into contributions from specific wiring features, then the spatial pattern encodes information about the mechanisms that generated it.

If spatial organization of synaptic inputs encodes something about dendritic computation, the arrangement of synapses should vary systematically with the electrotonic regime of the branch they innervate. Clustered inputs on isolated branches would support nonlinear integration; distributed inputs on compact branches would be consistent with linear summation (Branco et al., 2010; Mel, 1994). Alternatively, if spatial patterns reflect wiring logistics — which axons can physically reach which dendritic domains, and how many contacts they make — the coupling between regime and spatial organization would be explained by partner-level innervation structure rather than computational optimization (Voruganti, 2026).

A prior study on 12 manually curated MICrONS neurons reported preliminary evidence for the latter hypothesis: partner contact multiplicity mediated 69% of the coupling between electrotonic regime and synaptic compactness in excitatory neurons, with only 39% mediation in inhibitory neurons (Voruganti, 2026). However, the small sample limited statistical power and the precision of these estimates: one of three spatial metrics failed to reach significance, per-neuron consistency was low (2/12 individually significant), and the inhibitory subtype decomposition rested on $n = 2$ per subtype. Three critical questions remained unanswered: does partner-level mediation replicate at population scale? Does the $L/\lambda \rightarrow$ spatial coupling slope vary across neurons, or is it a fixed population property? And what explains the non-monotonic dose-response in pairwise compactness?

Here we address these questions by scaling the analysis to all 2,138 proofread neurons in the MICrONS volume (136,414 branches, ∼8.2 million synapses). We compute per-branch partner architecture features across the full population, fit random-slope mixed models that allow the coupling to vary per neuron, and resolve the compactness non-monotonicity through geometric diagnostics. We find that partner architecture mediates 25–53% of the coupling (exceeding the shallow 1–10% mediation by input balance), that excitatory and inhibitory neurons implement qualitatively different organizational logics, and that the compactness non-monotonicity is a normalization artifact rather than a biological transition zone. Critically, inhibitory neurons are not homogeneous: random-slope analysis reveals that the pooled inhibitory near-zero slope (-0.07) is an averaging artifact masking an inverted BC profile ($+1.81$) and BPC/NGC slopes steeper than the excitatory average (-2.76, -3.25), a pattern fully explained by subtype-specific dendritic domain preferences.

2 Results

2.1 Population-scale morphometric analysis

We extracted dendritic skeletons, synapse coordinates, and presynaptic partner identities for all 2,138 proofread neurons in the MICrONS mouse visual cortex volume (MICrONS Consortium, 2021; Turner et al., 2022) using the CAVEclient programmatic interface. Neurons spanned 11 cell types: seven excitatory subtypes (23P, $n = 518$; 4P, $n = 508$; 5P-ET, $n = 70$; 5P-IT, $n = 195$; 5P-NP, $n = 13$; 6P-CT, $n = 164$; 6P-IT, $n = 234$) and four inhibitory subtypes (basket cells [BC], $n = 213$; Martinotti cells [MC], $n = 114$; bipolar cells [BPC], $n = 85$; neurogliaform cells [NGC], $n = 24$). We identified 136,414 dendritic branches carrying a total of approximately 8.2 million synapses.

For each branch, we computed electrotonic length L/λ using cable theory with standard passive parameters ($R_M = 20{,}000\ \Omega \cdot \mathrm{cm}^2$, $R_I = 150\ \Omega \cdot \mathrm{cm}$) and branch-specific diameters derived from per-node radii in the EM reconstruction (Methods). The electrotonic length distribution was right-skewed (median $L/\lambda = 0.059$, IQR = 0.024–0.121), with 62.4% of branches below $L/\lambda = 0.1$ (electrotonically compact) and 4.5% above $L/\lambda = 0.5$ (strongly compartmentalized). For discretized analyses, we classified branches into three regimes using tertile thresholds on the L/λ distribution (regime 0: $L/\lambda < 0.037$; regime 1: $0.037 \leq L/\lambda < 0.157$; regime 2: $L/\lambda \geq 0.157$).

We quantified synaptic spatial organization using three complementary metrics, each z-scored against 1,000 uniform null draws per branch (Methods). The first two assess *local spacing structure* — whether neighboring synapses are evenly distributed along the branch: the **Clark-Evans ratio** (Clark and Evans, 1954) compares the mean nearest-neighbor distance to that expected under uniform placement, and the **Interval CV** measures the regularity of the gaps between consecutive synapses. The third assesses *global dispersion* — how spread out the entire synapse population is across the branch: **pairwise compactness** is the mean Euclidean distance between all synapse pairs. These two dimensions — local spacing and global dispersion — need not covary, a distinction that becomes important in Sections 2.5 and 2.7. Positive Clark-Evans z indicates more regular spacing than expected by chance; negative Interval CV z indicates more uniform inter-synapse intervals; positive compactness z indicates tighter clustering (Figure 1).

2.2 Continuous electrotonic length is associated with synaptic spatial organization

To avoid discretization artifacts, we first tested whether spatial organization covaries continuously with electrotonic length. We fit linear mixed-effects models with L/λ as a continuous predictor, controlling for synapse count and total branch length, with neuron identity as a random intercept (Methods).

Inhibitory neurons showed strong continuous coupling across all three metrics: Clark-Evans z increased with L/λ ($\beta = 7.65$, $p = 2.3 \times 10^{-77}$, $n = 22{,}808$ branches), Interval CV z decreased ($\beta = -17.5$, $p = 1.4 \times 10^{-90}$), and pairwise compactness z increased ($\beta = 11.7$, $p = 9.1 \times 10^{-117}$). The direction is consistent: longer, more isolated branches carry synapses that are more regularly spaced, with more uniform intervals, and that are on average more tightly clustered (Figure 2).

Excitatory neurons showed a more nuanced pattern. Interval CV z scaled strongly with L/λ ($\beta = -3.10$, $p = 1.0 \times 10^{-67}$, $n = 89{,}195$ branches), and Clark-Evans z showed a weaker but significant relationship ($\beta = -0.63$, $p = 2.5 \times 10^{-5}$, $n = 86{,}595$ branches). The negative sign of the excitatory Clark-Evans coefficient, opposite to the positive inhibitory coefficient, does not reflect a reversal of the biological direction: both cell types show more regular spacing on isolated branches. The sign difference arises because the mixed-effects model adjusts for covariates (synapse

count, branch length) that shift the baseline z-score distributions differently in excitatory and inhibitory populations (Figure 7). However, pairwise compactness showed no significant continuous relationship ($\beta = -0.07$, $p = 0.67$, $n = 89{,}200$ branches). This metric-specific divergence is resolved in Section 2.6, where random-slope models reveal that it reflects fundamentally different organizational logics rather than a simple absence of coupling.

2.3 Discretized regime analysis confirms graded coupling

To contextualize the continuous results and enable comparison with prior work, we classified branches into three electrotonic regimes using tertile thresholds and fit mixed-effects models with regime as a categorical predictor, controlling for synapse count, total branch length, and excitatory fraction of presynaptic inputs.

All three metrics showed large regime effects (Figure 3):

- Clark-Evans z: regime 2 coefficient $= 0.74$ (Cohen's $d = 0.60$; $n = 109{,}403$ branches)
- Interval CV z: regime 2 coefficient $= -1.70$ (Cohen's $d = -0.64$; $n = 112{,}202$)
- Pairwise compactness z: regime 2 coefficient $= 0.91$ (Cohen's $d = -0.07$; $n = 112{,}207$)

All regime effects were significant at $p < 10^{-300}$.

Dose-response analysis confirmed that the Clark-Evans and Interval CV effects are monotonically graded across regimes (Spearman $\rho = 0.25$ and $\rho = -0.37$, respectively; $p < 10^{-200}$ for both). The median Clark-Evans z progresses from -0.15 (regime 0) to 0.37 (regime 1) to 0.61 (regime 2), and the median Interval CV z from -0.05 to -0.78 to -1.14.

Sensitivity to threshold choice. We tested the robustness of the discretized result across eight threshold schemes: the primary tertile thresholds, biophysical thresholds ($L/\lambda = 0.1, 0.5$), and six quantile-based variants spanning 20th/80th to 40th/60th percentile splits. All eight converged for all three metrics (24/24 models), with Clark-Evans regime 2 coefficients ranging from 0.55 (biophysical) to 0.80 (20th/80th quantile) (Figure 1B). This stability demonstrates that the coupling is not an artifact of a particular threshold choice.

2.4 Cell-type dissociation: inhibitory coupling is uniformly stronger

The cell-type dissociation identified in prior work (Voruganti, 2026) replicates at population scale and extends to all four inhibitory subtypes tested (Figure 4).

For Clark-Evans z, inhibitory neurons show a regime 2 coefficient of 1.32 (Cohen's $d = 0.94$, $n = 22{,}808$) compared to 0.69 for excitatory neurons ($d = 0.60$, $n = 86{,}595$) — a 1.9-fold difference. This pattern holds uniformly across inhibitory subtypes:

- **Basket cells** ($n = 213$, 13,944 branches): regime 2 coefficient $= 1.68$, $d = 1.15$
- **Bipolar cells** ($n = 85$, 3,573 branches): regime 2 coefficient $= 1.16$, $d = 0.83$
- **Neurogliaform cells** ($n = 24$, 1,693 branches): regime 2 coefficient $= 1.08$, $d = 0.39$
- **Martinotti cells** ($n = 114$, 7,438 branches): regime 2 coefficient $= 0.79$, $d = 0.64$

Basket cells show the strongest coupling of any cell type, with effect sizes exceeding all excitatory subtypes (excitatory range: $d = 0.45$–1.14 across seven subtypes). Martinotti cells, despite being the weakest inhibitory subtype, still show coupling comparable to the excitatory average. Within excitatory neurons, layer 5 extratelencephalic-projecting neurons (5P-ET, $d = 1.00$) and layer 6 corticothalamic neurons (6P-CT, $d = 0.91$) show particularly strong effects, while layer 2/3 pyramidal cells (23P, $d = 0.50$) and layer 4 pyramidal cells (4P, $d = 0.45$) are weaker.

The Clark-Evans effect-size hierarchy across subtypes — BC > BPC > NGC > MC among inhibitory, 5P-ET > 6P-CT > 5P-IT > 6P-IT > 23P > 4P among excitatory — is consistent across fixed-effects metrics and suggests that the magnitude of aggregate regime coupling reflects cell-type-specific wiring architecture. Note that this hierarchy is metric-specific: the random-slope ICV analysis (Section 2.6) reveals a different ordering in which BC shows an inverted profile and 6P-CT shows the steepest excitatory slope, because interval-uniformity and nearest-neighbor-regularity track different aspects of the same underlying organization.

2.5 Partner architecture mediates regime coupling via two distinct mechanisms

The original 12-neuron study identified partner contact multiplicity as the primary mediator of regime coupling (69% mediation in excitatory neurons; Voruganti 2026). However, the population-scale replication (Section 2.8) found that excitatory/inhibitory input balance — the only partner-level feature tested at scale — mediates only 1–10% of the coupling, leaving a large mechanistic gap. We addressed this by computing four partner architecture features for each of 127,697 branches across all 2,138 neurons (Methods).

Per-branch partner architecture features. For each branch with matched presynaptic partners, we computed: (1) **unique partners**, the number of distinct presynaptic neurons; (2) **mean synapses per partner**, the ratio of total synapse count to unique partners (contact multiplicity); (3) **Gini coefficient of partner synapse counts**, measuring concentration inequality across partners; and (4) **fraction of multi-synaptic partners**, the proportion of partners making ≥ 2 contacts on the branch. Verification confirmed perfect correlation between the product of unique partners and mean synapses per partner and total synapse count ($r = 1.000$), validating the feature computation.

Individual mediator analysis. We tested each partner feature as a mediator using the coefficient-comparison framework (Methods). Throughout this paper, "mediation" refers to statistical decomposition of covariance: inclusion of a candidate mediator reduces the regime coefficient, indicating shared variance. This does not establish causal direction — partner architecture features could be causal mediators or correlated consequences of the same accessibility constraints that drive both regime assignment and spatial organization. The results revealed a striking dissociation between which features mediate which spatial dimensions (Table 2):

Partner diversity governs local spacing. For Clark-Evans z and Interval CV z — the two metrics that assess nearest-neighbor regularity and interval uniformity — **unique partners** was the strongest individual mediator in the overall population (23.0% and 22.9% mediation, respectively). Within excitatory neurons specifically, fraction of multi-synaptic partners was marginally stronger for Clark-Evans (23.7% vs. 22.7%), but the difference is within noise and both features reflect the same underlying partner diversity mechanism. This indicates that the number of distinct presynaptic neurons innervating a branch is the strongest predictor of how evenly synapses are spaced.

Contact multiplicity governs global dispersion. For pairwise compactness z — the metric that assesses how spread out synapse are across the branch — **mean synapses per partner** was the strongest individual mediator (39.3% overall). This indicates that how many contacts each partner makes predicts the spatial concentration of the overall synapse population.

All four partner features were individually significant mediators (all $p < 10^{-77}$), confirming that partner architecture operates through multiple correlated but non-redundant channels.

Combined mediation. When all four partner features were included simultaneously, combined mediation reached 28.9% for Clark-Evans z, 25.2% for Interval CV z, and 43.3% for pairwise compactness z in the overall population. This far exceeds the 1–10% mediation by input balance alone, establishing partner architecture as the dominant mediator of regime coupling at population scale.

Cell-type stratification reveals the inhibitory residual. The most important finding is the asymmetry between cell types. In excitatory neurons, partner architecture mediates 29–53% of the coupling (combined: 29.4% Clark-Evans, 31.3% Interval CV, 52.6% compactness). This is consistent with the original 12-neuron estimate of 69% mediation in excitatory neurons, converging from a larger population with more features.

In inhibitory neurons, by contrast, partner architecture mediates only 12–14% (combined: 14.0% Clark-Evans, 12.1% Interval CV, 12.6% compactness). For inhibitory compactness, the individual best mediator (mean synapses per partner) shows 15.0% mediation, but the combined model yields only 12.6% — a suppression effect indicating that some partner features act as negative confounders in the inhibitory context. This means that while inhibitory neurons show $\sim$2$\times$ stronger regime coupling than excitatory neurons (Section 2.4), substantially less of that coupling is explained by the wiring features we can measure. Something else — beyond partner diversity and contact multiplicity — is organizing inhibitory dendrites. We term this the **inhibitory residual**: the $\sim$86% of inhibitory coupling that cannot be attributed to partner-level innervation structure. Its identity is an open question addressed in the Discussion.

2.6 Excitatory and inhibitory neurons implement distinct organizational logics

The analyses above treat the $L/\lambda \rightarrow$ spatial coupling as a fixed population-level property, with neuron identity modeled only as a random intercept. This assumes every neuron has the same coupling slope. We tested this assumption by fitting random-slope mixed models that allow the $L/\lambda \rightarrow$ spatial metric slope to vary across neurons (Methods).

The coupling slope varies substantially across neurons. For all three metrics, the random-slope model was overwhelmingly preferred over the random-intercept-only model (likelihood ratio test: $\chi^2 = 1{,}724$, $p < 10^{-300}$ for Clark-Evans; $\chi^2 = 23{,}629$, $p < 10^{-300}$ for Interval CV; $\chi^2 = 1{,}021$, $p = 1.6 \times 10^{-222}$ for compactness; all with 2 df). The regime coupling is not a fixed property of the cortex — it varies per neuron.

Interval CV slopes reveal subtype-specific organizational logics. The most striking results emerged from the Interval CV slopes. We extracted per-neuron BLUP estimates for all 2,138 neurons and compared distributions across excitatory subtypes and inhibitory subtypes separately (Figure 8; Table 3).

Excitatory neurons show steep negative slopes across all subtypes: as L/λ increases, inter-synapse spacing becomes more uniform (population median slope $= -2.23$, $n = 1{,}702$ neurons).

Critically, the excitatory slope gradient by subtype is significant (Kruskal-Wallis $H = 61.1$, $p = 2.7 \times 10^{-11}$): L6-CT neurons show the steepest coupling (median $= -3.04$), substantially steeper than L5-ET neurons (median $= -1.84$, Mann-Whitney $p = 5.5 \times 10^{-11}$). The gradient follows circuit role rather than dendritic complexity: L6-CT neurons send corticothalamic feedback requiring precisely coordinated inputs, while L5-ET neurons integrate broadly across apical and basal compartments.

Inhibitory subtypes show divergent, bimodal regime-sensitivity profiles that average to near-zero at the population level. This population average (-0.07, $n = 436$) is an aggregation artifact that conceals fundamentally different subtype behaviors:

- **BC** (basket cells, PV-like): *positive* slope (median $= +1.81$) — more electrotonically isolated BC-postsynaptic branches have *less* regular spacing. This is consistent with BC operating outside its preferred proximal/compact territory on rare isolated branches, where inputs are sparser and less constrained.
- **MC** (Martinotti cells, SST-like): intermediate negative slope (median $= -0.80$).
- **BPC** (bipolar cells, VIP-like): median $= -2.76$, and **NGC** (neurogliaform cells): median $= -3.25$ — both steeper than the excitatory average.

BPC and NGC show stronger regime-coupling than the typical excitatory neuron, not weaker. The organizational diversity is therefore not excitatory versus inhibitory: it is subtype-specific, with BC defining one pole and BPC/NGC defining another. Figure 8 shows the full per-neuron slope distributions for all 11 subtypes.

Clark-Evans and compactness slopes. Clark-Evans slopes showed a modest but significant cell-type difference (excitatory median $= 0.91$, inhibitory median $= 1.14$; $p = 1.2 \times 10^{-4}$, $d = 0.16$). Compactness slopes did not differ between cell types (excitatory median $= 2.89$, inhibitory median $= 2.98$; $p = 0.21$, $d = -0.28$). The fact that global synapse dispersion (compactness) follows the same L/λ slope in both cell types, while local spacing regularity (Interval CV) does not, reinforces the distinction between the two spatial dimensions identified in the partner mediation analysis (Section 2.5).

Random slope variance. The random slope variance was largest for Interval CV ($\sigma^2_{\text{slope}} = 22.2$) — 17× larger than Clark-Evans ($\sigma^2_{\text{slope}} = 1.28$) and 4× larger than compactness ($\sigma^2_{\text{slope}} = 5.22$). This large variance is now interpretable: the Interval CV slope distribution is not unimodal but spans BC (positive slopes, $+1.81$) through BPC/NGC (strongly negative, -2.76 to -3.25), producing high population variance. The Interval CV slope is therefore not simply the "noisiest" metric — it is the metric that most sharply resolves subtype-specific organizational heterogeneity.

Reconciling fixed-effects, random-slope, and subtype results. Two apparent tensions in the results warrant explicit reconciliation.

First: Table 1 shows BC has the strongest Clark-Evans fixed-effects coupling of any cell type (regime 2 coefficient $= 1.68$, $d = 1.15$), yet Section 2.6 shows BC has a positive Interval CV random-slope ($+1.81$). These are not contradictory — they measure different spatial dimensions. Clark-Evans captures nearest-neighbor regularity, which increases with regime in BC (consistent with denser, more compact arrangements on compact branches). Interval CV captures inter-synapse *interval uniformity*, which becomes *less* regular on isolated branches in BC neurons. These can

simultaneously be true if BC contacts on compartmentalized branches are regularly spaced in absolute nearest-neighbor distance (Clark-Evans positive) but not in the gaps between consecutive synapses (ICV positive). Both metrics are internally consistent; the apparent contradiction dissolves when the two dimensions are distinguished.

Second: The fixed-effects model shows inhibitory regime coupling $\sim 2\times$ stronger than excitatory in aggregate, yet the subtype-stratified random-slope analysis shows the pooled inhibitory ICV slope is near-zero. These are also not contradictory. The strong aggregate inhibitory fixed-effects coupling is driven primarily by Clark-Evans (nearest-neighbor regularity) and compactness (global dispersion), not by Interval CV. The Interval CV channel — where the subtype heterogeneity is strongest — is the specific dimension on which the inhibitory population shows opposing within-subtype slopes that cancel in aggregate. The two analyses are complementary: fixed-effects models detect aggregate regime sensitivity, while random-slope analysis resolves the subtype-specific structure within that aggregate.

2.7 Pairwise compactness non-monotonicity is a normalization artifact

While Clark-Evans and Interval CV show clean monotonic dose-response patterns, pairwise compactness shows a non-monotonic relationship with electrotonic regime: regime 1 (intermediate) branches show *less* compact synapses than either regime 0 or regime 2. Geometric diagnostics (Figure 6) identify a regime $\times$ branch length interaction in the z-score normalization as the source: the compactness metric's sensitivity to branch length differs across regimes, producing apparent non-monotonicity in the normalized metric even when the underlying spatial organization varies monotonically with L/λ. We therefore present Clark-Evans and Interval CV as the primary evidence for graded regime coupling, with pairwise compactness treated as a secondary metric (see Supplementary for full diagnostics).

2.8 Partner-level validation of regime coupling

Two additional analyses confirm that the regime coupling reflects genuine partner-level wiring differences rather than a compositional artifact of input type.

Input balance does not explain regime coupling. The excitatory fraction of presynaptic inputs varies systematically across the dendritic tree, with distal (high-L/λ) branches typically receiving a different ratio of excitatory to inhibitory inputs than proximal branches. If this compositional gradient drives the spatial organization differences between regimes, controlling for excitatory fraction should eliminate the regime effect. Including excitatory fraction as a covariate reduced the regime 2 coefficient by 1.4% for Interval CV, 6.5% for Clark-Evans, and increased it by 7.0% for pairwise compactness (Table 4). The regime effect remained overwhelmingly significant in all cases ($p < 10^{-300}$). For inhibitory neurons, the mediation percentages were negligible to negative (-0.2% to -6.8%), indicating that controlling for input balance slightly *increases* the apparent regime effect. The coupling is therefore not a secondary consequence of input type composition.

Partner spatial spread scales with regime. To assess how individual presynaptic partners distribute their synapses across branches of different electrotonic regimes, we analyzed 178,103 partners making ≥ 3 synapses each. Partners innervating compartmentalized (regime 2) branches showed a mean pairwise spread of 91.3 μm, compared to 42.1 μm for partners on compact (regime 0) branches — a 2.17-fold difference (Figure 5). This confirms that the branch-level spatial orga-

nization patterns are accompanied by partner-level architectural differences: the same presynaptic neuron distributes its synapses more widely when contacting isolated dendritic domains.

2.9 Afferent subtype composition varies systematically with electrotonic regime

The inhibitory residual identified in Section 2.5 — the ∼86% of inhibitory coupling not attributable to partner architecture — could arise from cell-type-specific afferent targeting: if different presynaptic classes preferentially innervate branches of different electrotonic regimes, this compositional gradient could organize synaptic spacing through mechanisms orthogonal to generic partner diversity and contact multiplicity. We tested this by computing, for each dendritic branch of each inhibitory neuron, the fraction of typed presynaptic synapses originating from each inhibitory subtype (BC, MC, BPC, NGC) and from excitatory neurons (Methods).

BC afferents are strongly enriched on compact branches. Using a depth-preserving within-neuron shuffle null (1,000 permutations per neuron, preserving branch geometry and total per-class synapse counts within soma-distance tertiles; Methods), we found that basket cell afferents are strongly enriched on regime 0 (compact) branches: $\log_2$ enrichment $= 0.91$, $p < 0.001$ (permutation). At the neuron level, 88.9% of 434 inhibitory neurons showed higher BC fraction on compact than non-compact branches (median $\log_2$ ratio $= 0.99$, Hodges–Lehmann estimator $= 0.98$, Wilcoxon signed-rank $p = 4.0 \times 10^{-45}$, 95% bootstrap CI $= [0.92, 1.07]$). The effect was consistent across all four postsynaptic subtypes: BC postsynaptic neurons (97.7% positive, $n = 213$), MC (87.7%, $n = 114$), BPC (73.5%, $n = 83$), and NGC (70.8%, $n = 24$).

Within-neuron targeting tests confirmed that this is not a between-neuron confound: 190 of 213 individual BC neurons showed significant within-neuron BC afferent targeting ($p < 0.05$, chi-squared test), with 208/213 showing the expected directional pattern (BC→low regime, MC→high regime).

MC targeting is real but modest. Martinotti cell afferents showed a statistically reliable but substantially weaker enrichment on compartmentalized branches (neuron-level median $\log_2$ ratio $= 0.26$, Hodges–Lehmann $= 0.30$, Wilcoxon $p = 2.3 \times 10^{-19}$, 67.4% positive). The BC effect is approximately 4× larger than the MC effect, consistent with the stronger compartment specificity of perisomatic-targeting interneurons compared to distal-targeting Martinotti cells (Markram et al., 2004; Schneider-Mizell et al., 2023).

BC targeting concentrates in proximal dendritic compartments. Stratifying by soma distance (per-neuron median split), the BC targeting effect was 16× stronger in proximal compartments (mixed-model regime 2 coefficient $= -0.30$, $p = 6.2 \times 10^{-147}$) than in distal compartments (coefficient $= -0.019$, $p = 0.036$). This proximal concentration is consistent with the perisomatic targeting biology of basket cells (Schneider-Mizell et al., 2023) but resolves finer electrotonic structure than compartment-level analyses alone: within proximal dendrites, BC afferents still preferentially target electrotonically compact branches.

Mixed-effects models confirmed the targeting pattern at population level: frac_BC_afferent $\sim C(\text{regime})$, with regime 2 coefficient $= -0.28$ ($p < 10^{-300}$, $n = 25{,}902$ branches, 436 neurons), and frac_MC_afferent: regime 2 coefficient $= +0.073$ ($p = 4.1 \times 10^{-39}$).

2.10 Afferent identity and partner architecture jointly explain 18–26% of inhibitory coupling

To determine whether afferent subtype composition adds explanatory power beyond the partner architecture features identified in Section 2.5, we fit a four-model comparison for each spatial metric (Table 5):

- **M0** (baseline): $z_{\text{metric}} \sim C(\text{regime}) + \text{synapse_count} + \text{total_length} + (1|\text{neuron})$
- **M1** (+ composition): M0 + frac_inh + frac_BC + frac_MC
- **M2** (+ partner architecture): M0 + unique_partners + mean_syn/partner + Gini + frac_multisynaptic
- **M3** (+ both): M0 + all composition and partner architecture covariates

For all three spatial metrics, M3 > M1 and M3 > M2: the combined model consistently explained more variance than either pathway alone. Combined mediation (M3) reached 23.3% for Clark-Evans z, 18.1% for Interval CV z, and 25.5% for pairwise compactness z — up from 14.2%, 13.3%, and 15.4% for partner architecture alone (M2). All bootstrap 95% CIs excluded zero (200 cluster-bootstrap iterations by neuron; Table 5).

Two partially independent pathways. The sequential decomposition reveals two pathways through which afferent identity organizes synaptic spacing:

The overlap pathway (4–6%): afferent composition acts *through* partner architecture. BC targeting predicts which presynaptic neurons arrive, which predicts contact multiplicity and partner diversity, which predicts spacing. This was confirmed directly: BC fraction significantly predicted all four partner architecture features in mixed models (Link 1 mediation: 13–39%; BC fraction coefficients $p < 10^{-18}$ for all four features). BC-rich branches have fewer unique partners ($\beta = -2.20$, $p = 6.9 \times 10^{-19}$), higher contact multiplicity ($\beta = +0.064$, $p = 6.4 \times 10^{-83}$), higher partner concentration ($\beta = +0.039$, $p = 2.0 \times 10^{-83}$), and more multisynaptic contacts ($\beta = +0.043$, $p = 1.5 \times 10^{-85}$) — precisely the partner architecture profile that drives compact spacing.

The chain incremental pathway (5–10%): afferent composition explains regime variance in spatial metrics *beyond* what partner architecture captures. This placement effect indicates that different presynaptic classes impose different spacing rules conditional on where they synapse, a hypothesis tested directly in Section 2.11.

Subtype-specific chain patterns. Stratifying by postsynaptic subtype confirmed that the chain incremental is concentrated in BC (5–11%) and MC (5–9%) postsynaptic neurons, while BPC neurons are dominated by partner architecture alone (M2 = 21–67%, chain incremental = 1–5%). This dissociation indicates that the mechanisms differ by postsynaptic subtype: BC and MC postsynaptic neurons benefit from afferent-specific placement rules, while BPC organization is largely explained by generic partner architecture.

2.11 Presynaptic subtype imposes distinct within-axon placement rules

The chain incremental (Section 2.10) demonstrates that afferent identity affects spatial organization beyond partner architecture, but does not reveal whether this reflects a genuine placement effect or a residual confound. We tested this directly with an axon-conditioned placement analysis: for multisynaptic typed partners (same presynaptic neuron with ≥ 2 synapses on the same branch), we

computed the mean pairwise 3D distance between their synapses and tested whether electrotonic regime predicts this within-partner placement metric after conditioning on contact multiplicity and branch length (Methods).

BC contacts are placed more compactly on compact branches. Among 7,478 BC partner-branch records (407 neurons), the mean within-partner pairwise distance increased monotonically with electrotonic regime: 2.5 μm (regime 0, n = 2,560), 3.3 μm (regime 1, n = 2,069), and 4.4 μm (regime 2, n = 2,849) — a 1.76-fold expansion from compact to compartmentalized. In a mixed-effects model controlling for contact count (k) and branch length, the regime 2 coefficient was +0.95 ($p = 1.1 \times 10^{-107}$; Table 6).

The effect survives exact multiplicity matching. To eliminate any residual multiplicity confound, we analyzed partners at each contact count separately. The regime gradient persisted at every k value tested: $k = 2$ (regime 2 coefficient = 0.93, $p = 1.8 \times 10^{-67}$, n = 5,578), $k = 3$ (= 1.07, $p = 4.7 \times 10^{-39}$, n = 1,274), and $k = 4$ (= 1.00, $p = 4.1 \times 10^{-18}$, n = 407). This eliminates the possibility that the placement effect is an artifact of differential contact multiplicity across regimes. (Note: equal-k subsets are nested within the "All" sample; they partition it by contact count rather than constituting independent replications.)

Neuron-level confirmation. At the neuron level, 81.0% of 300 neurons with BC partners in both regime 0 and regime 2 showed wider BC contact placement on compartmentalized branches (median $\log_2$(R2/R0) = 0.90, Hodges–Lehmann = 0.97, Wilcoxon $p = 1.2 \times 10^{-31}$, 95% CI = $[0.76, 1.08]$).

MC contacts spread more broadly than BC contacts on compartmentalized branches. MC partners showed a qualitatively similar but quantitatively larger expansion: median within-partner distance increased from 2.7 μm (regime 0) to 8.1 μm (regime 2) — a 3.0-fold expansion compared to BC's 1.76-fold. An interaction model (regime $\times$ presynaptic class, controlling for k and branch length) confirmed that the regime 2 $\times$ MC interaction was significant ($\beta = +0.24$, $p = 5.2 \times 10^{-4}$), meaning MC contacts spread significantly more on compartmentalized branches than BC contacts do, conditional on being there.

This subtype-specific placement difference — BC tight, MC broad — aligns with the known axonal morphology of these classes: basket cell axons form dense perisomatic bouton clusters, while Martinotti cell axons arborize more broadly in distal dendritic domains (Markram et al., 2004). The axon-conditioned analysis shows that this morphological difference manifests at the level of individual branch electrotonic structure, not just at the coarse compartment level.

Mechanistic interpretation. Together, Sections 2.9–2.11 trace a multi-stage pathway that partially resolves the inhibitory residual:

electrotonic regime $\rightarrow$ *afferent subtype composition* $\rightarrow$ *partner architecture* $\rightarrow$ *within-axon placement* $\rightarrow$ *synaptic spacing*

Each link is independently supported. The arrival effect (Section 2.9) predicts which presynaptic classes innervate a branch. The partner architecture effect (Section 2.5) predicts the diversity and multiplicity of contacts. The placement effect (this section) predicts where individual axons position their boutons. These three mechanisms are partially independent: their combined mediation (M3

= 18–26%) exceeds either pathway alone, and the placement effect survives exact multiplicity conditioning. The pathway explains ∼25% of the inhibitory residual, leaving ∼75% attributable to wiring features not resolved at partner-level granularity — most likely sub-partner-scale placement programs of the kind characterized in the B1/B2 analysis (see Discussion, Section 3.2).

2.12 Robustness checks

Volume-edge truncation diagnostic. Dendritic branches near the boundary of the MICrONS acquisition volume may be truncated, potentially compressing synapse positions and artificially inflating spacing regularity metrics on those branches. To test whether such truncation drives the regime coupling results, we computed the minimum distance from each branch centroid to the inferred volume boundary (estimated from the 1st–99th percentile of all SWC node positions: X = 499–1378 μm, Y = 357–1041 μm, Z = 632–1079 μm). Regime 0, 1, and 2 branches had nearly identical median boundary distances (R0: 98.4 μm, R1: 99.0 μm, R2: 96.2 μm), indicating that compact and compartmentalized branches are not differentially distributed near the volume boundary. We then re-ran the regime coupling models excluding branches within 0, 10, 25, and 50 μm of the boundary. Across all populations (all neurons, excitatory, inhibitory) and all three spatial metrics, the regime 2 coefficient shifted by ≤10.9% at the most aggressive exclusion threshold (50 μm), with the largest shift in the conservative direction (excluding boundary branches *reduces* the coupling slightly, ruling out inflation). We conclude that volume-edge truncation does not drive the observed regime coupling.

Presynaptic coverage ceiling. To assess whether the 43% typed-partner coverage ceiling (among $k \geq 3$ partners) could substantially bias the afferent composition results, we performed a systematic CAVE query for all available annotation tables in the `minnie65_public` datastack beyond the primary metamodel table (v661). The supplementary query (mtype table v661_v2, perisomatic ultrastructure entries) added only 87 new subtype assignments across 3.5 million unique presynaptic root IDs, establishing that 93% of presynaptic fragments lack any cell-body annotation in CAVE. The 43% coverage ceiling is therefore a practical constraint of the EM reconstruction, not a query limitation, and the typed-subset analyses operate on the analyzable population.

Within-axon placement: axon identity versus context. A potential confound for the placement analysis (Section 2.11) is that the observed BC-compact-branch association could reflect stable between-axon differences (some BC axons are inherently more compact placers) rather than context-sensitive responses to postsynaptic regime. We tested this by decomposing the residual variance in within-axon placement into between-axon and within-axon-across-regime components. Intraclass correlations across branches of the same axon were low (BC ICC = 0.032, MC ICC = 0.049), indicating that most within-axon variance is context-driven rather than axon-identity-driven. Consistent with this, same-axon cross-regime analyses showed that 79% of BC axons with contacts in multiple regimes spread *more* on compartmentalized branches (i.e., the placement difference is directional but context-dependent; Spearman $\rho = -0.006$, $p = 0.90$ for axon-mean vs. regime difference), and the arrival–placement correlation (BC $\rho = -0.267$, $p = 2.7 \times 10^{-11}$) shows that axons arriving preferentially on compact branches also place more compactly. This correlation is more than a robustness check: it demonstrates that domain-targeting tendency and placement precision are coordinated at the single-axon level. Axons with the strongest compact-branch preference implement the tightest compact-branch placement, suggesting that arrival strategy and bouton positioning reflect a coherent, axon-intrinsic relationship with postsynaptic branch type. These results confirm the placement effect is a contextual response rather than a fixed axonal property.

Electrotonic length versus raw geometry. A fundamental question is whether L/λ captures a biophysically specific property of dendritic branches, or whether it is merely a reparameterization of geometric descriptors (branch length, diameter, soma distance, branch order) that would predict spacing equally well. We tested this by fitting the cable formula $L/\lambda = L/\sqrt{R_M d/(4R_I)}$ to the four geometry features: the regression achieves $R^2 = 1.000$, confirming that L/λ is a deterministic function of length and diameter (as expected analytically). We then tested whether the specific $L/\sqrt{d}$ combination — the biophysically relevant ratio — predicts spacing metrics with incremental power beyond length and diameter entered as independent linear covariates. Across both cell types and all three metrics, L/λ significantly improved model fit after geometry control (LRT $p < 10^{-14}$ for all six tests; inhibitory ICV: $\chi^2 = 62.3$, $p = 2.9 \times 10^{-15}$; excitatory ICV: $\chi^2 = 1276$, $p < 10^{-279}$). This indicates the *ratio* of length to the square root of diameter — not each dimension independently — is the predictive quantity, consistent with the biophysical interpretation that electrotonic isolation, rather than physical geometry alone, organizes synaptic spacing. Notably, BC-afferent fraction was not significantly predicted by L/λ beyond geometry ($p = 0.08$), suggesting that BC targeting may operate primarily through positional cues (soma distance, branch order) rather than the biophysical regime per se.

B1 versus B2: origin of inhibitory spacing regularity. We tested whether the near-uniform inhibitory spacing arises from a postsynaptic scaffold that repels all inhibitory contacts regardless of their presynaptic identity (B1: gephyrin-mediated exclusion), or from within-axon self-avoidance where each presynaptic axon distributes its own contacts to cover territory (B2: presynaptic tiling). We computed inhibitory-synapse nearest-neighbour distances on the same branch, stratified by same-axon versus different-axon pairs. Under B1, same-axon and cross-axon NN distances should be comparable (the scaffold excludes all nearby contacts). Under B2, same-axon NN should be substantially larger than cross-axon NN (within-axon contacts are spaced out; contacts from different axons can land anywhere). Across 13,568 branches from 327 inhibitory neurons, same-axon NN distances (median = 2.34 μm) were 5.6× larger than cross-axon NN distances (median = 0.42 μm), and cross-axon NN distances were 2.6× smaller than a within-branch shuffle null (median = 1.65 μm). This strongly supports B2: inhibitory spacing regularity is a presynaptic property of individual axons, while different axons can and do land in close proximity. An additional observation is that same-axon contacts spread further apart on compartmentalized branches (regime 2 same-axon NN: 3.17 μm vs. regime 0: 1.02 μm), consistent with the contextual placement result (Section 2.12) — axons adapt their spacing to the postsynaptic branch length available.

Density-regularity relationship and the residual noise floor. Two further diagnostics clarify the nature of the inhibitory residual. First, we tested whether ICV regularity scales with synapse density — a key discriminator between hard-core repulsion (higher density forces tighter constraints → more regular) and coverage tiling (regularity independent of density). In inhibitory neurons, ICV z-score became significantly more negative with increasing synapse density ($\beta = -0.78$, $p < 10^{-130}$, LME controlling for branch length and neuron identity). Inhibitory ICV z improved monotonically from -0.62 (lowest density quintile, median 0.76 syn/μm) to -1.25 (fourth quintile, 2.34 syn/μm), with a partial reversal at the highest densities (-0.73 at fifth quintile, 3.47 syn/μm). The reversal at high densities is consistent with a tiling program that saturates near a coverage threshold: once contacts are sufficiently numerous to tile available territory, additional contacts are placed without further spatial constraint, reducing apparent regularity. This density-dependence is consistent with the B2 (within-axon tiling) interpretation: each axon places contacts to cover branch territory, so more contacts on the same branch are more constrained spatially. The same pattern holds in ex-

citatory neurons ($\beta = -0.95$, $p \approx 0$), confirming it reflects a general placement constraint rather than an inhibitory-specific feature.

Second, we estimated whether the ∼75% post-regime residual variance is measurement noise or real biology. Using an analytic approximation of ICV sampling variance ($\approx 1/[2(k-1)]$ for k synapses per branch), the estimated noise floor is 0.0189 ICV z-score units^2, compared to a post-regime residual variance of 7.76 units^2 — a signal-to-noise ratio of 345×. The inhibitory residual is definitively not measurement noise. It reflects real within-neuron, within-regime variation in spacing that is not captured by electrotonic regime, partner architecture, or afferent subtype identity — the signature of individual axonal tiling programs operating at sub-connectomic resolution.

3 Discussion

3.1 Two wiring mechanisms, not one

The partner architecture mediation analysis resolves the mechanistic shallowness of the initial population-scale replication. Where input balance (excitatory fraction) mediated only 1–10% of the regime coupling, partner architecture features mediate 25–53%, establishing partner-level innervation structure as the dominant mechanistic layer.

The dissociation between which features mediate which spatial dimensions is not redundancy — it supports two separable organizational channels. *Partner diversity* (the number of unique presynaptic neurons contacting a branch) is the strongest predictor of local spacing regularity: branches innervated by more distinct partners show more evenly spaced synapses. This is consistent with a geometric packing constraint — more partners competing for postsynaptic territory produce more regular spacing through mutual exclusion (Kasthuri et al., 2015). *Contact multiplicity* (how many synapses each partner makes per branch) is the strongest predictor of global synapse dispersion: branches where individual partners make repeated contacts show more concentrated synapse populations. This is consistent with the "multi-contact" wiring rule, in which axons that have already established one contact tend to make additional contacts nearby (Druckmann et al., 2014).

These two mechanisms can coexist because they operate at different spatial scales: partner diversity predicts the macro-scale regularity of the synapse population, while contact multiplicity predicts the micro-scale clustering of individual partner contributions. The mediation hierarchy — partner diversity for Clark-Evans and Interval CV, contact multiplicity for compactness — maps directly onto the local-spacing versus global-dispersion distinction identified in the metric design (Section 2.1).

3.2 Decomposing the inhibitory residual

The most consequential finding from the mediation analysis is the asymmetry between cell types. Partner architecture mediates 29–53% of excitatory coupling but only 12–14% of inhibitory coupling, despite inhibitory neurons showing ∼2× stronger regime effects. The afferent decomposition (Sections 2.9–2.11) partially decomposes this asymmetry, reducing the inhibitory residual from ∼86% to ∼75% through a multi-stage mechanistic chain.

The chain has three independently validated links. First, electrotonic regime predicts afferent subtype composition: basket cell afferents concentrate on compact branches with a neuron-level effect that is consistent across 89% of inhibitory neurons. This targeting is not merely a depth effect — it persists within proximal dendritic compartments, resolving finer electrotonic structure than compartment-level analyses (Schneider-Mizell et al., 2023). Second, afferent subtype composition predicts partner architecture: BC-rich branches show the precise partner profile — fewer unique

partners, higher contact multiplicity, greater partner concentration — that drives compact spacing. Third, presynaptic subtype identity imposes distinct within-axon placement rules that survive exact multiplicity conditioning. This third link is the most novel: it demonstrates that the spatial organization of inhibitory synapses reflects not only *who* arrives and *how often*, but also *how* each presynaptic class positions its boutons once it has arrived.

The ∼75% that remains unexplained is interpretable in light of the B1/B2 test (Section 2.12). The result — that inhibitory spacing regularity is a within-axon property (B2: presynaptic tiling) rather than a postsynaptic scaffold property (B1) — identifies the residual as primarily presynaptic in origin: each inhibitory axon tiles its contacts across the branch, and this tiling is context-dependent (adapting to branch length rather than being a fixed inter-bouton rule). The ∼75% is therefore not a gap in the mechanistic chain but a presynaptic program operating at a finer scale than connectomic partner-level features can resolve. A parsimonious generative account of the full inhibitory picture is a two-stage rule: (i) *domain gating*, in which each inhibitory subtype has a coarse target-domain policy consistent with known PV perisomatic and SST distal-dendritic biology (Schneider-Mizell et al., 2023); and (ii) *territorial coverage*, in which once inside its domain, each axon tiles its contacts across available dendritic territory using a within-axon self-avoidance program (Section 2.12). Under this model, the ∼75% residual is not "unknown biology" but the presynaptic tiling program operating at finer spatial scale than partner-level connectomic features can resolve.

The domain-gating step has a structural circuit interpretation: BC and MC afferents implement anatomically distinct inhibitory functions. BC regulates somatic output by concentrating on electrotonically compact (proximal) dendritic domains, where it can efficiently shunt excitatory drive regardless of its dendritic location. MC selectively gates dendritic computation by targeting compartmentalized (distal) branches capable of local nonlinear integration, where its contacts can veto dendritic events at the branch level (Larkum et al., 2009). This functional specialization is structurally supported by BC's regime-0 enrichment in 89% of neurons and MC's regime-2 enrichment (Section 2.9). Direct functional demonstration would require paired structural–functional recording at the dendritic level.

Arrival versus placement. The multi-stage pathway distinguishes two conceptually distinct mechanisms. The *arrival* effect (Section 2.9) determines which presynaptic classes are present on a branch — BC afferents concentrate on compact branches, MC afferents on compartmentalized branches. The *placement* effect (Section 2.11) determines how each class distributes its boutons once it has arrived — BC contacts are placed more tightly on compact branches even after conditioning on how many contacts they make. These are separable: the four-model comparison (Section 2.10) shows that arrival acts partly through partner architecture (the overlap pathway, 4–6%) and partly through an independent channel (the chain incremental, 5–10%). The axon-conditioned analysis confirms that the placement channel survives exact multiplicity matching, ruling out the possibility that arrival and placement are merely two descriptions of the same multiplicity-driven process.

3.3 Subtype-specific organizational logics for synaptic spacing

The subtype-stratified random-slope analysis reveals that neither "excitatory" nor "inhibitory" is a homogeneous category with respect to synaptic spacing organization. The dissociation previously framed as a qualitative excitatory-versus-inhibitory divide is more precisely a spectrum of subtype-specific regime-coupling profiles.

For excitatory neurons, the coupling between electrotonic isolation and spacing regularity is real and consistent across subtypes (all medians negative), but the magnitude varies significantly by

subtype. A mixed-effects interaction model ($L/\lambda\times$ subtype, with L2/3 [23P] as reference) confirms that all deeper-layer excitatory subtypes show significantly steeper slopes than L2/3 neurons (4P: $\Delta = -0.046$, $p = 3.2 \times 10^{-7}$; 6P-CT: $\Delta = -0.230$, $p = 9.6 \times 10^{-59}$; 5P-IT: $\Delta = -0.140$, $p = 2.2 \times 10^{-33}$). Among excitatory subtypes, L6-CT shows the strongest coupling (median $= -3.04$) and L5-ET the shallowest (median $= -1.84$). A potential confound is that L6-CT branches have a significantly higher L/λ distribution than L5-ET branches (median 0.14 vs. 0.08; Mann-Whitney $p = 10^{-83}$; 59% vs. 45% of branches above $L/\lambda = 0.1$). The formal interaction model partially controls for this by testing whether slope differs across subtypes at any given continuous L/λ value (6P-CT interaction coefficient $\Delta = -0.230$, $p = 9.6 \times 10^{-59}$), suggesting the subtype effect is not purely a distribution artifact. However, the contribution of dendritic morphology differences between L6-CT (relatively compact, ascending axons) and L5-ET (complex apical dendrites, broad dendritic arbor) to the slope difference cannot be fully separated from circuit-role effects in the current data. The gradient is consistent with a model in which circuit role shapes the degree of spatial sorting — L6-CT's corticothalamic feedback function may require precisely coordinated input timing — but this interpretation is supported by the pattern, not established by a causal test.

For inhibitory neurons, the pooled median (-0.07) is an averaging artifact masking opposing subtype behaviors. The formal interaction model (vs. BC reference) confirms BPC significantly differs from BC ($\Delta = -0.324$, $p = 4.1 \times 10^{-12}$) and MC significantly differs from BC ($\Delta = -0.135$, $p = 2.5 \times 10^{-5}$); NGC is directionally consistent but underpowered ($n = 24$, $p = 0.23$).

BC neurons show an *inverted* coupling profile (median $= +1.81$): on the rare electrotonically isolated branches that BC contacts, spacing is less regular, not more. This is not an outlier-driven result — 63.8% of individual BC neurons show positive per-neuron slopes, and the pattern survives under coarser molecular taxonomy (PV-like, SST-like, VIP-like, NGC-like; Kruskal-Wallis $H = 189.4$, $p = 7.0 \times 10^{-40}$; PV-like vs. VIP-like Mann-Whitney $p = 4.7 \times 10^{-9}$). This is the parsimonious signature of a domain-specialist cell type operating outside its preferred territory: BC afferent fraction drops sharply with electrotonic regime (regime 2 coefficient $= -0.28$, $p < 10^{-300}$; Section 2.9), meaning high-L/λ branches genuinely receive sparse BC contacts. On those branches, BC inputs arrive as rare visitors without the territorial coverage constraints that govern placement in preferred proximal domains. The inverted slope is therefore the spatial footprint of domain targeting, not a distinct spacing algorithm. This interpretation is consistent with BC's known proximal targeting specialization (Schneider-Mizell et al., 2023) and not a demonstrated mechanism.

BPC and NGC, by contrast, show the strongest negative slopes in the entire dataset (-2.76 and -3.25), steeper than the average excitatory neuron. BPC neurons, associated with VIP-positive disinhibitory circuits that gate local cortical computation (Pfeffer et al., 2013), may require precisely coordinated inputs whose spatial organization matters for the timing of disinhibitory gates. NGC neurons operate through GABA volume transmission and spillover — mechanisms that may be particularly effective on electrotonically isolated branches where local GABA accumulation is less rapidly dissipated. The regime-sensitivity of NGC input organization is consistent with this volume-transmission context dependence, though the direct mechanism remains to be established. MC occupies an intermediate position (median $= -0.80$), consistent with its distal-dendritic targeting role (Schneider-Mizell et al., 2023).

The organizational diversity is therefore not a binary excitatory/inhibitory division but a continuum structured by subtype identity and circuit role. The common thread across all subtypes with strongly negative slopes is that they are likely to receive inputs whose spatial coordination matters for the downstream computation: excitatory neurons receiving inputs whose clustering or regularity affects dendritic integration, and non-perisomatic inhibitory neurons (BPC, NGC) whose inputs may require similar spatial precision for local inhibitory circuit function. BC, whose function

is fast, broad inhibition of somatic output rather than spatially precise dendritic control, is the one cell type where regime-matched spacing regularity is not expected — and is not observed.

Together, these profiles are consistent with a domain-matching principle: each inhibitory subtype's spacing regularity peaks in the electrotonic compartment that matches its functional target zone. BC (perisomatic target) degrades on isolated branches precisely because it seldom operates there; MC (distal-dendritic target) shows a negative slope, consistent with regularity increasing toward compartmentalized domains (Schneider-Mizell et al., 2023); BPC and NGC, operating across the dendritic tree, show the steepest magnitude slopes, indicating uniform spatial-coordination pressure regardless of regime. This principle generates a testable prediction: the MC ICV slope should shift toward more negative values if the effective target compartment of MC shifts (e.g., across cortical areas or projection classes with different apical-dendritic extent), and toward less negative values in neurons receiving unusually sparse distal MC input.

Structure, function, and plasticity history. We note that the per-neuron ICV coupling slope does not predict orientation selectivity (OSI) or response reliability in the 45.8% of proofread neurons with MICrONS functional coregistration ($n = 845$ matched excitatory neurons; Spearman $\rho = -0.003$, $p = 0.94$ for ICV slope vs. OSI; see Limitation 8 for full results). This is expected under a plasticity-history interpretation: if excitatory spatial organization reflects the cumulative record of synaptic modification over the neuron's lifetime, structural organization encodes what the neuron *has learned*, while OSI reflects the neuron's *current stimulus-driven response state*. These need not be correlated if ongoing experience continuously recalibrates tuning while the structural trace of past plasticity accumulates more slowly. A direct test of this interpretation would require dark-rearing or monocular deprivation paradigms, where excitatory slopes should be shallower in deprived circuits if the structural organization reflects activity-dependent input redistribution. We state this as a prediction, not a conclusion.

The fact that compactness slopes do not differ substantially between cell types ($p = 0.21$) reinforces that this subtype specificity is confined to local interval regularity — the dimension most directly reflecting active biological regulation. Global synapse dispersion (compactness) responds similarly across all subtypes, consistent with physical constraints (axonal reach, branch length) governing that dimension universally.

3.4 Resolving the compactness non-monotonicity

The non-monotonic dose-response in pairwise compactness is resolved as a normalization artifact (Section 2.7; Supplementary). This does not invalidate pairwise compactness as a metric — it remains informative about global synapse dispersion within a regime. But cross-regime comparisons of compactness z-scores should be interpreted cautiously when the regimes differ substantially in branch length. Clark-Evans and Interval CV are less susceptible to this artifact because nearest-neighbor and interval metrics are less sensitive to branch geometry than pairwise distance metrics.

3.5 Relation to the original 12-neuron study

The present results both confirm and extend the original finding of 69% partner-mediated coupling in excitatory neurons (Voruganti, 2026). At population scale, partner architecture mediates 29–53% of excitatory coupling, depending on the metric. The lower bound is expected: the original study used a single feature (mean synapses per partner) with a more aggressive matching procedure on 12 manually curated neurons, while the present study uses four features across 1,702 excitatory neurons with automated matching. The convergence on the same order of magnitude — tens of

percent, not single digits — validates the original mechanism while providing more precise and generalizable estimates.

The inhibitory mediation finding (12–14%) is novel. The original study estimated 39% mediation for inhibitory neurons, but with $n = 2$, this estimate had wide uncertainty. The population-scale result establishes that inhibitory coupling is genuinely less mediated by partner architecture than excitatory coupling, confirming the cell-type asymmetry as a robust biological property rather than a small-sample artifact.

3.6 Limitations

Several limitations should be noted. First, all neurons are from a single MICrONS cortical volume (mouse visual cortex area VISp). While the within-volume replication is comprehensive (all 2,138 proofread neurons), cross-dataset generalization — to other cortical areas, other species (e.g., H01 human cortex), or subcortical structures — remains untested.

Second, electrotonic lengths are computed from EM-measured diameters, which are subject to tissue shrinkage during electron microscopy preparation. The absolute L/λ values are therefore approximate. However, the sensitivity analysis demonstrates that the coupling is invariant to threshold choice across a wide range, and the continuous analysis uses relative L/λ rankings rather than absolute values.

Third, the passive cable parameters (R_M, R_I) are fixed at standard values. In vivo, active conductances, synaptic conductance loads, and neuromodulation alter the effective electrotonic length of a branch dynamically. Our analysis captures the static structural regime, not the instantaneous functional regime. Importantly, varying R_M or R_I scales all L/λ values by a common factor ($L/\lambda \propto \sqrt{R_M/R_I}$) without changing the rank ordering of branches. Because the primary regime classification uses tertile thresholds on the L/λ distribution (not biophysically absolute thresholds), the regime assignments are invariant to uniform scaling of cable parameters. The continuous analysis (Section 2.2) uses L/λ as a relative predictor whose direction and significance are also invariant to scaling. This property, combined with the robustness to eight alternative threshold schemes demonstrated in Section 2.3, means that the results are not sensitive to the specific choice of cable parameters within the range of commonly used values.

Fourth, the afferent subtype labels (BC, MC, BPC, NGC) are morphological EM proxy categories, not molecular identity classes. The correspondence with molecular types (PV, SST, VIP, 5HT3aR) is approximate (Schneider-Mizell et al., 2023): our claims about presynaptic identity should be interpreted as "consistent with" known molecular targeting biology rather than direct demonstrations of molecular identity.

Fifth, of all presynaptic synapses, only ∼18% have resolved cell type annotations; among multi-synapse partners ($k \geq 3$), coverage reaches ∼43%. The afferent composition analyses therefore operate on the typed subset, which may not be fully representative of all inputs. Coverage is regime-asymmetric (proximal branches have lower subtyped coverage than distal), but this asymmetry works *against* finding BC enrichment on compact branches — making the observed effect conservative.

Sixth, the combined mechanistic chain (afferent identity + partner architecture) explains 18–26% of inhibitory coupling, leaving ∼75% unexplained. Candidate mediators not tested include cell-autonomous spacing mechanisms, developmental timing, axonal geometry constraints, and molecular signals inaccessible to connectomic data.

Seventh, the random-slope model AIC/BIC values were numerically undefined (NaN) due to the large sample size and model complexity, though the likelihood ratio tests and parameter estimates were valid. The model comparison should be interpreted through the LRT statistics rather than

information criteria.

Eighth, we pre-registered and tested whether the per-neuron ICV slope predicts orientation selectivity (OSI) or response reliability (cc_abs) in the 45.8% of proofread neurons with functional coregistration ($n = 845$ matched excitatory neurons; Spearman $\rho = -0.003$, $p = 0.94$ for ICV slope vs. OSI; $\rho = -0.048$, $p = 0.16$ for cc_abs). All three pre-registered hypotheses were null. As interpreted in Discussion Section 3.3, this null is consistent with a plasticity-history interpretation in which structural organization encodes past synaptic modification while OSI reflects current tuning state; these dissociate under ongoing experience. Whether the null instead reflects insufficient coregistration power (45.8% match rate), measurement noise, or a genuine absence of structure-function coupling cannot be determined from the present data alone.

Finally, we do not link spatial organization to neural computation more broadly. Whether the spatial patterns have functional consequences — for example, whether the excitatory spacing-regularity coupling produces different patterns of dendritic integration than the inhibitory cell-autonomous spacing — requires combined structural and functional data with higher coregistration completeness (Jia et al., 2010).

4 Methods

4.1 Data acquisition

We queried the MICrONS mouse visual cortex EM volume (MICrONS Consortium, 2021; Turner et al., 2022) using CAVEclient (v8.0.1) with the `minnie65_public` datastack. Cell type annotations were obtained from the AIBS metamodel cell types table (v661); proofreading status from the proofreading status table. All 2,178 proofread neurons were downloaded; 2,138 passed quality control (see below). Dendritic skeletons with per-node radii were obtained via the skeleton service. Postsynaptic synapse positions were queried from the `synapses_pni_2` table. Presynaptic partner cell types were resolved by batch-querying the cell type table for unique presynaptic root IDs. Voxel coordinates were converted to nanometers using the volume resolution $[4, 4, 40]$ nm/voxel.

4.2 Quality control

Neurons were excluded if: (a) the skeleton failed to parse or had zero dendrite branches, (b) fewer than 20 synapses were recovered, (c) fewer than 80% of synapses snapped successfully to the skeleton (snap distance threshold 50 μm), or (d) fewer than 50% of presynaptic partners had resolved cell type annotations. Of 2,178 downloaded neurons, 2,138 passed all criteria (98.2%).

4.3 Branch morphometry and electrotonic length

Dendritic skeletons were parsed into branches at bifurcation points. For each branch, we computed physical length (sum of inter-node Euclidean distances in nm) and mean diameter (twice the mean of per-node radii from the EM reconstruction, in nm). Electrotonic length was computed as:

$$L/\lambda = \frac{L}{\sqrt{R_M \cdot d/(4R_I)}}$$

where L is physical length (cm), d is diameter (cm), $R_M = 20{,}000\ \Omega \cdot \mathrm{cm}^2$, and $R_I = 150\ \Omega \cdot \mathrm{cm}$ (Rall, 1967; Koch, 1999).

4.4 Regime classification

Branches were classified into three electrotonic regimes using tertile thresholds on the population L/λ distribution: regime 0 ($L/\lambda < 0.037$, compact), regime 1 ($0.037 \leq L/\lambda < 0.157$, intermediate), and regime 2 ($L/\lambda \geq 0.157$, compartmentalized). For sensitivity analysis, we additionally tested biophysical thresholds ($L/\lambda = 0.1, 0.5$) and six quantile-based thresholds spanning 20th/80th to 40th/60th percentile splits.

4.5 Spatial organization metrics

For each branch with ≥ 3 synapses, we computed three spatial metrics. The first two capture local spacing structure (are neighboring synapses evenly distributed?), while the third captures global dispersion (how spread out is the entire synapse population?).

Clark-Evans ratio (Clark and Evans, 1954): the ratio of mean observed nearest-neighbor distance to the expected mean under a uniform distribution on the branch path. Values > 1 indicate regular spacing; < 1 indicates clustering.

Interval CV: the coefficient of variation of inter-synapse intervals along the branch path. Lower values indicate more uniform spacing.

Pairwise compactness: the mean Euclidean distance between all pairs of synapses on the branch. Lower values indicate tighter spatial concentration.

4.6 Null model and z-scoring

Each metric was z-scored against a dendrite-constrained null model. For each branch, we drew 1,000 independent samples of n uniformly distributed points along the branch path (where n is the observed synapse count) and computed the metric for each null draw. The z-score is $(x_{\text{obs}} - \mu_{\text{null}})/\sigma_{\text{null}}$.

All spatial metrics are computed on the full synapse population per branch, including both excitatory and inhibitory presynaptic synapses. Synapse-type-specific metrics (e.g., ICV computed on inhibitory inputs only) were not computed because the majority of branches contain insufficient numbers of labeled synapses of a single type to support reliable metric estimation (typical labeled inhibitory synapse count per branch is 1–3; the minimum for ICV estimation is 3). The z-scoring against a branch-specific null partially mitigates the synapse-type mixture effect by conditioning on the observed total synapse count, but reviewers should note that the computed metrics reflect the spatial organization of the full synaptic population rather than excitatory or inhibitory inputs independently.

4.7 Statistical testing

We fit linear mixed-effects models using the `statsmodels` MixedLM implementation in Python:

$$z_{\text{metric}} \sim \text{regime} + \text{synapse_count} + \text{total_length} + \text{exc_fraction} + (1|\text{neuron})$$

where regime is treated as a categorical variable (reference: regime 0), synapse count and total branch length are continuous covariates, excitatory fraction is the proportion of presynaptic partners classified as excitatory, and neuron identity is a random intercept. For continuous analyses, regime is replaced by continuous L/λ. Significance was assessed via Wald z-tests on fixed-effect coefficients.

Effect sizes were computed as Cohen's d between regime 0 and regime 2 z-score distributions. Dose-response monotonicity was assessed via Spearman rank correlation between regime label and metric z-score.

4.8 Partner architecture features

For each branch, presynaptic partner identities were assigned by matching synapse coordinates to presynaptic partner CSV files via KDTree nearest-neighbor lookup. For each branch with ≥ 3 snapped synapses, we computed: (1) **unique partners** — the count of distinct presynaptic root IDs; (2) **mean synapses per partner** — total synapse count divided by unique partners; (3) **Gini coefficient** — the Gini index of the per-partner synapse count distribution, measuring inequality (0 = all partners contribute equally, 1 = one partner dominates); (4) **fraction multi-synaptic** — the proportion of partners making ≥ 2 contacts on the branch. Features were computed for 127,697 branches across 2,138 neurons.

4.9 Mediation analysis

Mediation was tested using the coefficient-comparison approach (Baron and Kenny, 1986). For each mediator, we compared:

- **Reduced model:** $z_{\text{metric}} \sim \text{C(regime)} + \text{synapse_count} + \text{total_length} + (1|\text{neuron})$
- **Full model:** reduced model + mediator variable(s)

Mediation percentage was computed as $100 \times (1 - |\beta_{\text{full}}|/|\beta_{\text{reduced}}|)$, where β is the regime 2 coefficient. Positive values indicate that the mediator absorbs part of the regime effect; negative values indicate suppression (the mediator reveals additional regime effect). We note that mediation here is a statistical decomposition of covariance, not a demonstration of causal direction: partner architecture features could be causal mediators or correlated consequences of the same accessibility constraints that drive both regime assignment and spatial organization.

4.10 Random-slope mixed models

To test whether the $L/\lambda \rightarrow$ spatial coupling varies across neurons, we fit random-slope models:

$$z_{\text{metric}} \sim L/\lambda + \text{synapse_count} + \text{total_length} + (L/\lambda \,|\, \text{neuron})$$

This allows each neuron to have its own L/λ slope (BLUP estimate) while estimating a population-average fixed effect. The random-slope model was compared to the random-intercept-only model via likelihood ratio test (LRT; χ^2 with 2 df for the slope variance and intercept-slope covariance). Per-neuron slopes were extracted as the sum of the fixed effect and the neuron-specific random slope. Excitatory vs. inhibitory slope distributions were compared using Mann-Whitney U tests and Cohen's d.

4.11 Compactness geometry diagnostics

The non-monotonic pairwise compactness dose-response was investigated through: (1) comparison of branch length and synapse density distributions across regimes; (2) OLS residual analysis, regressing compactness z on synapse count and total branch length, then testing whether residuals still differ by regime; (3) interaction model, adding regime $\times$ branch length terms to test whether the compactness–length relationship has different slopes across regimes.

4.12 Partner spatial spread

For each presynaptic partner making ≥ 3 synapses onto a postsynaptic neuron, we computed the mean 3D Euclidean distance between all pairs of their synapse positions. Partners were classified by the dominant electrotonic regime of the branches they innervate (the regime containing the plurality of their synapses).

4.13 Afferent subtype composition

For each dendritic branch of each inhibitory neuron, we matched presynaptic synapses to partner identity using KDTree nearest-neighbor lookup against presynaptic partner CSV files. Each synapse was assigned an annotation level: L1 (exact subtype from CAVE annotation), L2 (subtype from catalog backfill), L3 (broad class only: excitatory/inhibitory), or L4 (unknown). Afferent fractions (frac_BC, frac_MC, frac_BPC, frac_NGC) were computed using L1+L2 synapses; broad fractions (frac_exc, frac_inh) used L1+L2+L3.

4.14 Depth-preserving shuffle null

For each postsynaptic neuron, we assigned each synapse to a soma-distance tertile bin (proximal/medial/distal, computed per neuron). Within each bin, presynaptic subtype labels were randomly permuted 1,000 times, preserving branch geometry, synapse count per branch, total per-class synapse counts, and depth structure. The observed afferent fraction per regime was compared to the null distribution. Sensitivity analysis with 2 and 4 depth bins confirmed that results were not bin-dependent.

4.15 Four-model mediation comparison

For each spatial metric, we fit four nested mixed-effects models on inhibitory branches: M0 (baseline: regime + synapse count + branch length + random intercept for neuron), M1 (M0 + afferent fractions), M2 (M0 + partner architecture features), and M3 (M0 + both). Mediation percentage was computed as $100 \times (1 - |\beta_{\text{full}}|/|\beta_{\text{reduced}}|)$. The chain incremental (= M3 − M2) measures the contribution of afferent composition beyond partner architecture; the overlap (= M1 + M2 − M3) measures the shared pathway. Bootstrap 95% CIs were computed via 200 cluster-bootstrap iterations (resampling neurons with replacement, fitting OLS).

4.16 Axon-conditioned placement analysis

For each inhibitory neuron, we identified multisynaptic typed partners: presynaptic neurons with ≥ 2 synapses of known subtype on the same branch. For each partner-branch pair, we computed the mean 3D pairwise Euclidean distance between all synapse pairs. The placement metric was modeled as: $\log(1+\text{mean_pairwise_dist}) \sim C(\text{regime}) + k_{\text{synapses}} + \text{total_length} + (1|\text{neuron})$. The equal-$k$ analysis stratified by exact contact count ($k = 2, 3, 4$) and refit within each stratum. The interaction model tested $C(\text{regime}) \times C(\text{pre_class})$ with BC as reference.

4.17 Software

All analyses used Python 3.9 with NumPy, pandas, SciPy, and statsmodels. Skeleton parsing and synapse snapping used custom code. Figures were generated with Matplotlib using the Wong (2011) colorblind-safe palette.

Table 1: Summary of neuron counts, branch counts, and Clark-Evans regime 2 coupling coefficients by cell type. Cohen's d computed between regime 0 and regime 2 Clark-Evans z-score distributions.

Class	Subtype	Neurons	Branches	Regime 2 coef.	Cohen's d	p
Inhibitory	BC	213	13,944	1.68	1.15	$< 10^{-290}$
	BPC	85	3,573	1.16	0.83	$< 10^{-37}$
	NGC	24	1,693	1.08	0.39	$< 10^{-12}$
	MC	114	7,438	0.79	0.64	$< 10^{-38}$
Excitatory	5P-ET	70	9,330	1.34	1.00	$< 10^{-126}$
	5P-NP	13	389	1.19	1.14	$< 10^{-5}$
	6P-CT	164	7,111	1.06	0.91	$< 10^{-80}$
	5P-IT	195	11,090	0.81	0.72	$< 10^{-64}$
	6P-IT	234	9,353	0.79	0.72	$< 10^{-51}$
	23P	518	48,328	0.68	0.50	$< 10^{-149}$
	4P	508	24,165	0.61	0.45	$< 10^{-79}$
All inhibitory		**436**	**26,648**	**1.32**	**0.94**	$< 10^{-300}$
All excitatory		**1,702**	**109,766**	**0.69**	**0.60**	$< 10^{-300}$

Table 2: Partner architecture mediation of regime coupling. Mediation %: proportion of the regime 2 coefficient absorbed by partner features. "Combined" includes all four features simultaneously. For comparison, excitatory fraction mediates only 1–10%.

Population	Metric	Individual best	Med. %	Combined med. %	n
Overall	Clark-Evans z	unique_partners	23.0%	28.9%	121,826
	Interval CV z	unique_partners	22.9%	25.2%	127,678
	Compactness z	mean_syn/partner	39.3%	43.3%	127,686
Excitatory	Clark-Evans z	frac_multisynaptic	23.7%	29.4%	97,000
	Interval CV z	unique_partners	29.0%	31.3%	102,387
	Compactness z	mean_syn/partner	48.5%	52.6%	102,394
Inhibitory	Clark-Evans z	unique_partners	13.2%	14.0%	24,826
	Interval CV z	unique_partners	11.2%	12.1%	25,291
	Compactness z	mean_syn/partner	15.0%	12.6%	25,292

Table 3: Random-slope mixed model results. Fixed slope: population-average L/λ coefficient. Slope variance: between-neuron variance in the L/λ slope. LRT compares random-slope to random-intercept-only model (2 df). Per-neuron slopes extracted as BLUPs.

Metric	Fixed slope	p	Slope var.	LRT χ^2	Exc. med.	Inh. med.	MW p
CE z	0.97	2.8×10^{-11}	1.28	1,724	0.91	1.14	1.2×10^{-4}
ICV z	−1.22	5.4×10^{-8}	22.24	23,629	−2.23	−0.07	1.1×10^{-24}
Comp. z	2.80	1.0×10^{-58}	5.22	1,021	2.89	2.98	0.21

Table 4: Mediation of regime coupling by excitatory/inhibitory input balance. Negative mediation values indicate that controlling for input balance increases the apparent regime effect.

Population	Metric	β_{reduced}	β_{full}	Mediation %
Overall	Clark-Evans z	0.794	0.743	6.5%
	Interval CV z	−1.720	−1.697	1.4%
	Compactness z	0.851	0.910	−7.0%
Excitatory	Clark-Evans z	0.765	0.690	9.8%
	Interval CV z	−1.359	−1.302	4.2%
	Compactness z	0.763	0.773	−1.2%
Inhibitory	Clark-Evans z	1.312	1.315	−0.2%
	Interval CV z	−2.673	−2.825	−5.7%
	Compactness z	1.694	1.810	−6.8%

Table 5: Four-model comparison for inhibitory regime coupling. M0: baseline (regime + controls). M1: + afferent composition. M2: + partner architecture. M3: + both. Mediation %: reduction in regime 2 coefficient relative to M0. Chain incremental = M3 − M2 (composition beyond partner architecture). Overlap = M1 + M2 − M3 (shared pathway). Bootstrap 95% CIs from 200 cluster-bootstrap iterations (in brackets).

Metric	M1 %	M2 %	M3 %	Chain	Overlap	n
Clark-Evans z	13.9	14.2	23.3 [17.7, 23.0]	9.1 [6.5, 10.1]	4.7 [2.8, 4.5]	24,826
Interval CV z	8.6	13.3	18.1 [13.2, 19.4]	4.9 [3.1, 6.2]	3.8 [2.4, 3.6]	25,291
Compactness z	16.6	15.4	25.5 [20.3, 25.9]	10.2 [8.8, 12.8]	6.4 [4.8, 6.8]	25,292

Table 6: Axon-conditioned placement analysis. Within-partner mean pairwise 3D distance for multisynaptic typed partners ($k \geq 2$ synapses on the same branch). Regime coefficients from mixed models controlling for k and branch length (log-transformed distance). Equal-k rows show the regime 2 coefficient when holding contact count exactly constant.

Class	Subset	n	Neurons	R0 (μm)	R2 (μm)	R2 coef	p
BC	All	7,478	407	2.5	4.4	+0.95	1.1×10^{-107}
	$k = 2$	5,578	403	2.1	3.7	+0.93	1.8×10^{-67}
	$k = 3$	1,274	304	3.0	5.7	+1.07	4.7×10^{-39}
	$k = 4$	407	191	3.6	7.4	+1.00	4.1×10^{-18}
MC	All	4,468	411	2.7	8.1	+0.79	8.3×10^{-20}
	$k = 2$	3,525	399	2.4	6.3	+0.77	1.2×10^{-13}
	$k = 3$	663	264	3.0	13.8	+0.88	6.8×10^{-7}
Exc	All	12,072	414	2.5	3.9	+0.46	1.7×10^{-21}

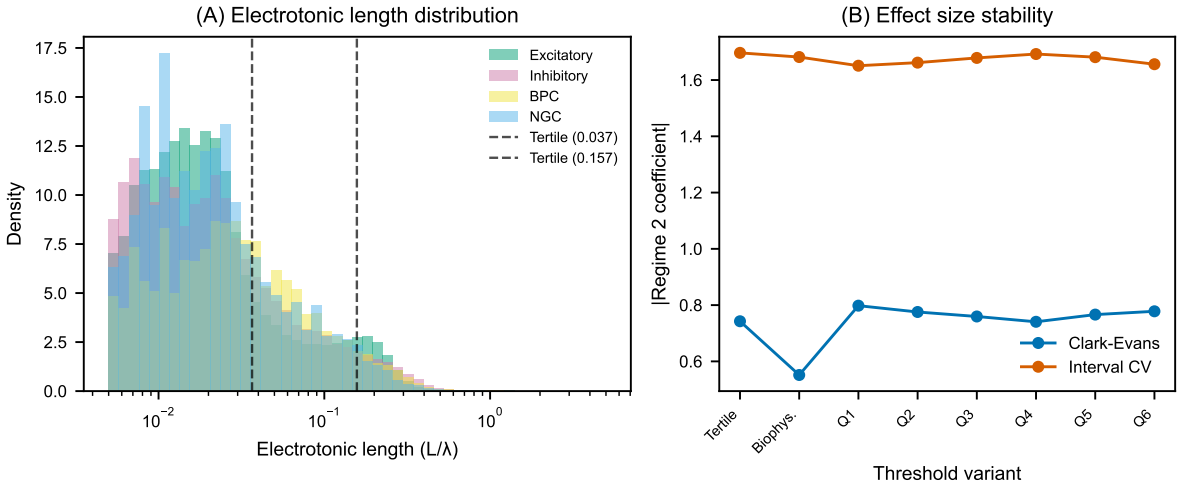

Figure 1: **Electrotonic length distribution and sensitivity analysis.** (A) Distribution of branch electrotonic lengths (L/λ) across 136,414 branches from 2,138 neurons. Vertical dashed lines indicate tertile thresholds. (B) Sensitivity of the regime 2 coefficient to threshold choice across eight discretization schemes for three spatial metrics. All 24 models converge, with Clark-Evans coefficients ranging from 0.55 to 0.80.

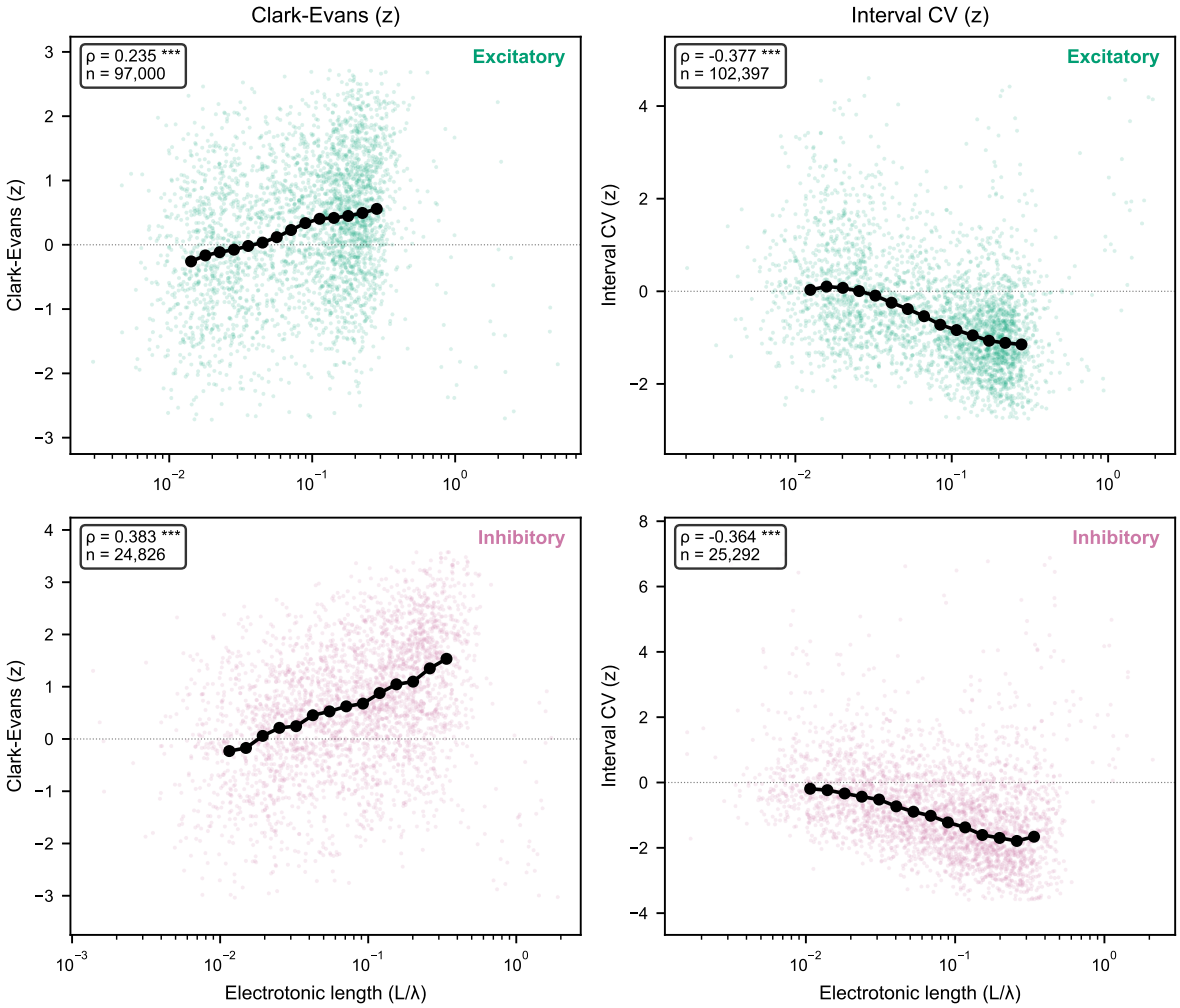

Figure 2: **Synaptic spatial organization tracks with continuous electrotonic length.** Scatter plots of L/λ vs. spatial metric z-scores for excitatory (top) and inhibitory (bottom) neurons, with binned medians (orange) computed in the 5th–95th percentile core (≥ 20 points per bin). Left: Clark-Evans z. Right: Interval CV z. Inhibitory neurons show steep, monotonic trends ($\beta = 7.65$, $p < 10^{-77}$ and $\beta = -17.5$, $p < 10^{-90}$, respectively). Excitatory neurons show significant but weaker continuous relationships.

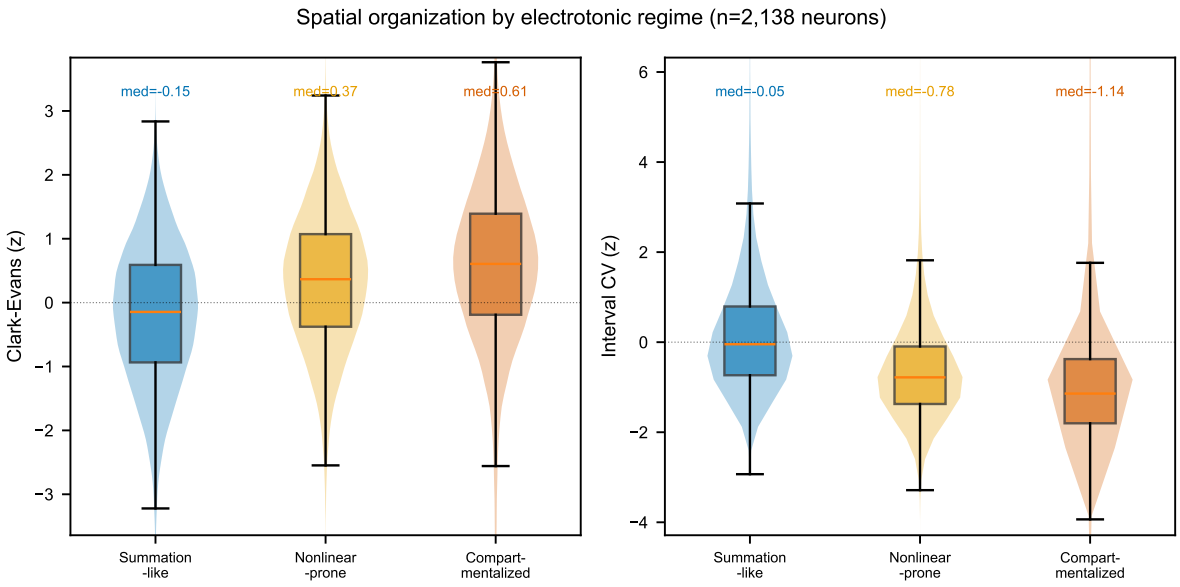

Figure 3: **Discretized regime analysis.** Violin plots of spatial metric z-scores stratified by electrotonic regime (0: compact, 1: intermediate, 2: compartmentalized) for 2,138 neurons. Clark-Evans z and Interval CV z show monotonic dose-response. Pairwise compactness z shows non-monotonic behavior explained by a regime $\times$ branch length normalization interaction (Section 2.7). All regime 2 vs. regime 0 comparisons: $p < 10^{-300}$.

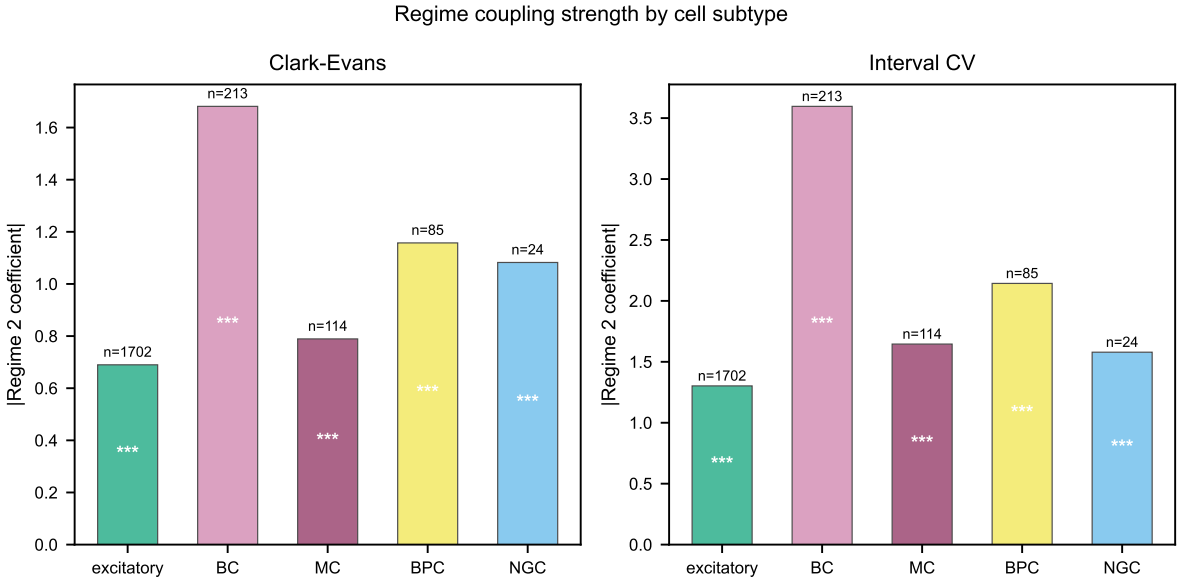

Figure 4: **Cell-type-specific coupling strength.** Clark-Evans regime 2 coefficients and Cohen's d effect sizes for 11 cell types. All four inhibitory subtypes (blue shades) show stronger coupling than the excitatory aggregate (orange). Basket cells are the strongest ($d = 1.15$); Martinotti cells are the weakest inhibitory subtype ($d = 0.64$) but still exceed the excitatory mean ($d = 0.60$). Error bars: 95% confidence intervals from mixed-effects models.

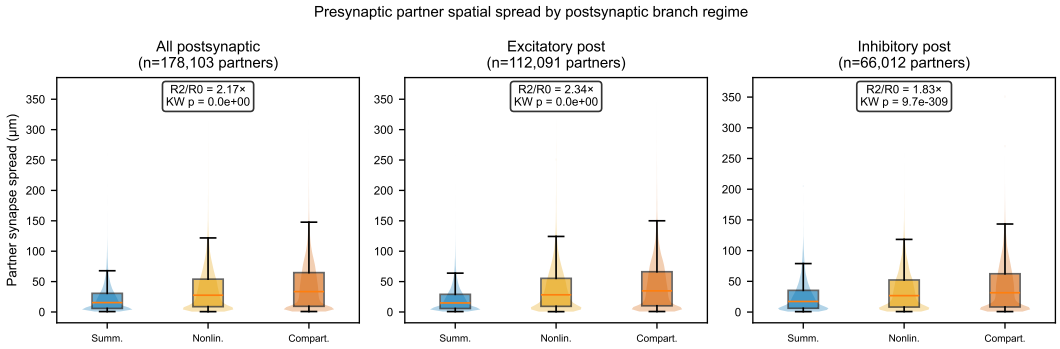

Figure 5: **Partner spatial spread by electrotonic regime.** Mean 3D pairwise distance of synapses from individual presynaptic partners ($k \geq 3$ synapses), classified by the dominant regime of innervated branches. Partners on regime 2 branches spread their synapses 2.17-fold wider than partners on regime 0 branches ($n = 178{,}103$ partners).

Supplementary Figures

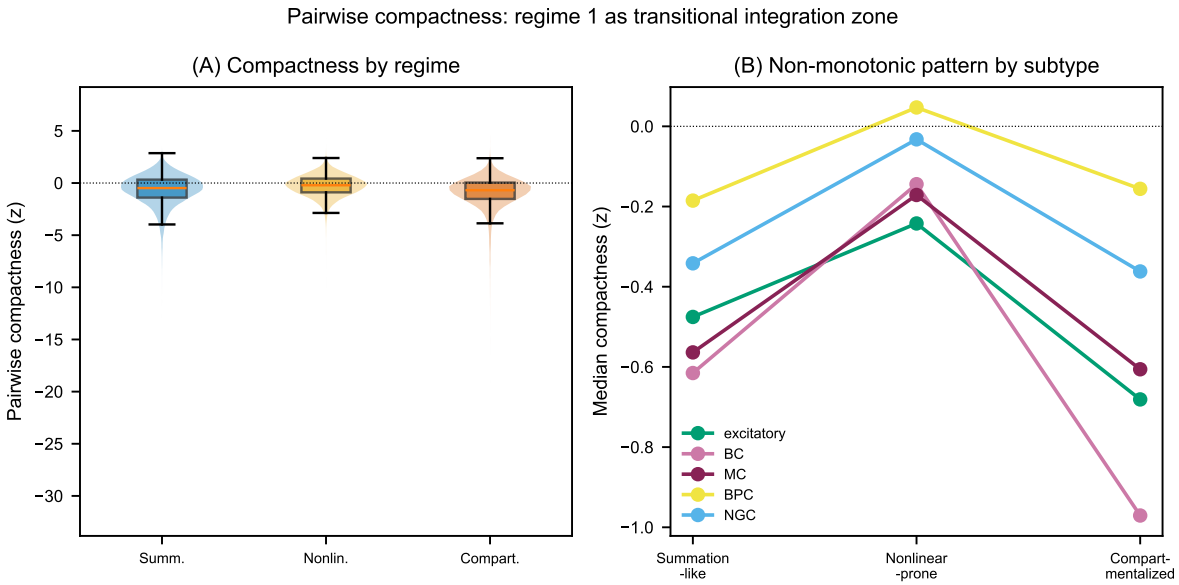

Figure 6: **Figure S1. Pairwise compactness non-monotonicity: normalization artifact across cell types.** Dose-response plots of median pairwise compactness z-score by electrotonic regime, stratified by cell type. The regime 1 reversal is consistent across excitatory and all four inhibitory subtypes, ruling out a subtype-specific artifact. Geometric diagnostics (Section 2.7) demonstrate that this non-monotonicity arises from a regime $\times$ branch length interaction in the z-score normalization: the compactness metric's sensitivity to branch length differs across regimes, producing apparent non-monotonicity in the normalized metric even when the underlying spatial organization varies monotonically with L/λ.

Compactness non-monotonicity: full diagnostics

The non-monotonicity in pairwise compactness z-scores arises from a regime $\times$ branch length interaction. By construction, branches in different regimes have different physical lengths: regime 0 median $= 11.6$ μm, regime 1 $= 42.3$ μm, regime 2 $= 104.7$ μm (Kruskal-Wallis $H = 110{,}679$, $p < 10^{-300}$). After regressing out synapse count and total branch length from pairwise compactness z, the residuals still differ by regime (regime 0 mean $= -0.37$, regime 1 $= +0.21$, regime 2 $= +0.11$; $H = 3{,}926$, $p < 10^{-300}$). A regime $\times$ branch length interaction model revealed significant interactions for both regime 1 ($\beta = 5.2 \times 10^{-5}$, $p < 10^{-300}$) and regime 2 ($\beta = 4.8 \times 10^{-5}$, $p < 10^{-300}$). The compactness z-score normalizes by the expected spread of uniformly distributed synapses, which scales with branch length. When this relationship has different slopes across regimes — as it does, because intermediate-length branches have specific geometric properties — the z-scored metric can show non-monotonic behavior even if the underlying spatial organization is monotonic.

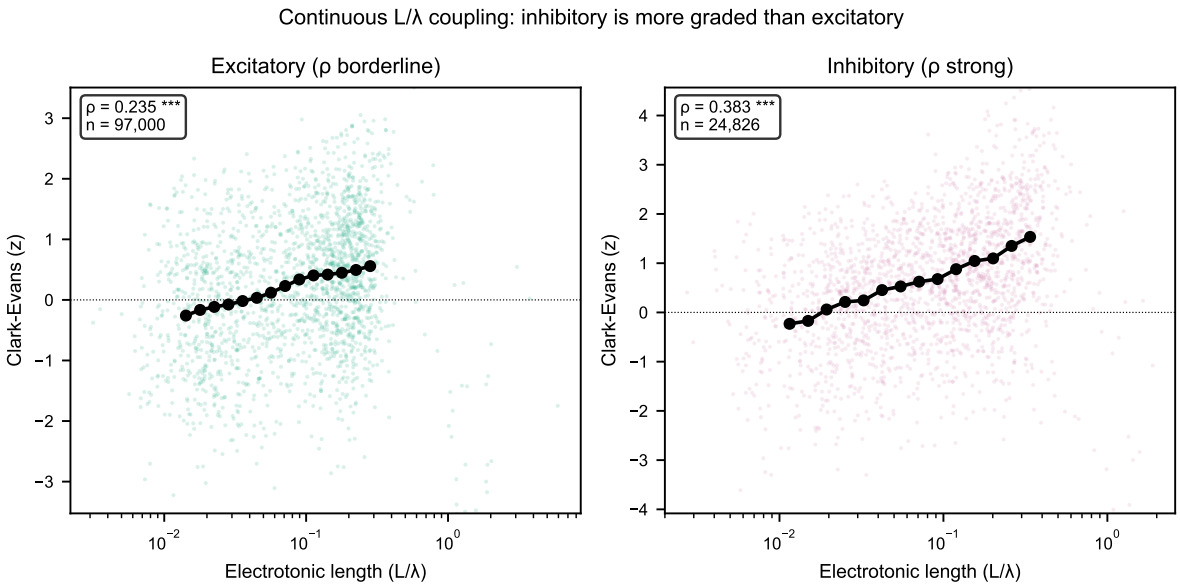

Figure 7: **Figure S2. Excitatory vs. inhibitory continuous coupling.** Direct comparison of L/λ vs. spatial metric z-scores between excitatory and inhibitory neurons. Inhibitory neurons show steeper slopes across all metrics. For pairwise compactness, the excitatory continuous relationship is null ($p = 0.67$) while the inhibitory relationship is strong ($p < 10^{-117}$), demonstrating that the metric-specific divergence is cell-type-dependent.

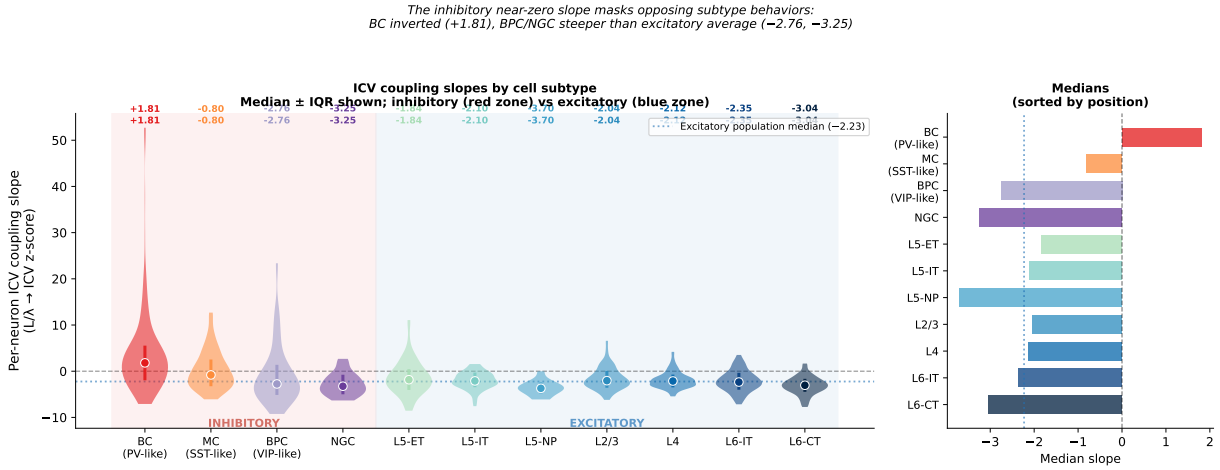

Figure 8: **Figure 6. Per-neuron ICV coupling slopes by cell subtype. (A)** Violin plots of per-neuron ICV coupling slopes (BLUP estimates from random-slope mixed model; $L/\lambda \rightarrow$ ICV z-score) for all 11 inhibitory and excitatory subtypes. Points show medians; vertical bars show IQR. Inhibitory subtypes (red zone, left) and excitatory subtypes (blue zone, right). The inhibitory pooled median (-0.07, dotted horizontal line in B) is an averaging artifact: BC shows an inverted profile ($+1.81$) while BPC (-2.76) and NGC (-3.25) show steeper coupling than the excitatory average (-2.23, blue dashed line). **(B)** Median slopes ranked by value, showing the full range from BC (positive) to NGC (most steeply negative). Organizational diversity is structured by subtype and circuit role, not a binary excitatory/inhibitory divide.

References

Baron, R. M. and Kenny, D. A. (1986). The moderator-mediator variable distinction in social psychological research. *Journal of Personality and Social Psychology*, 51(6):1173–1182.

Branco, T., Clark, B. A., and Häusser, M. (2010). Dendritic discrimination of temporal input sequences in cortical neurons. *Science*, 329(5999):1671–1675.

Clark, P. J. and Evans, F. C. (1954). Distance to nearest neighbor as a measure of spatial relationships in populations. *Ecology*, 35(4):445–453.

Druckmann, S., Feng, L., Lee, B., Yook, C., Zhao, T., Magee, J. C., and Kim, J. (2014). Structured synaptic connectivity between hippocampal regions. *Neuron*, 81(3):629–640.

Jia, H., Rochefort, N. L., Chen, X., and Konnerth, A. (2010). Dendritic organization of sensory input to cortical neurons in vivo. *Nature*, 464(7293):1307–1312.

Kasthuri, N., Hayworth, K. J., Berger, D. R., Schalek, R. L., Conchello, J. A., Knowles-Barley, S., Lee, D., Vázquez-Reina, A., Kaynig, V., Jones, T. R., et al. (2015). Saturated reconstruction of a volume of neocortex. *Cell*, 162(3):648–661.

Koch, C. (1999). *Biophysics of Computation: Information Processing in Single Neurons*. Oxford University Press.

Larkum, M. E., Nevian, T., Sandler, M., Polsky, A., and Schiller, J. (2009). Synaptic integration in tuft dendrites of layer 5 pyramidal neurons: a new unifying principle. *Science*, 325(5941):756–760.

London, M. and Häusser, M. (2005). Dendritic computation. *Annual Review of Neuroscience*, 28:503–532.

Markram, H., Toledo-Rodriguez, M., Wang, Y., Gupta, A., Silberberg, G., and Wu, C. (2004). Interneurons of the neocortical inhibitory system. *Nature Reviews Neuroscience*, 5(10):793–807.

Mel, B. W. (1994). Information processing in dendritic trees. *Neural Computation*, 6(6):1031–1085.

MICrONS Consortium (2021). Functional connectomics spanning multiple areas of mouse visual cortex. *bioRxiv*.

Pfeffer, C. K., Xue, M., He, M., Huang, Z. J., and Scanziani, M. (2013). Inhibition of inhibition in visual cortex: the logic of connections between molecularly distinct interneurons. *Nature Neuroscience*, 16(8):1068–1076.

Poirazi, P., Brannon, T., and Mel, B. W. (2003). Pyramidal neuron as two-layer neural network. *Neuron*, 37(6):989–999.

Polsky, A., Mel, B. W., and Schiller, J. (2004). Computational subunits in thin dendrites of pyramidal cells. *Nature Neuroscience*, 7(6):621–627.

Rall, W. (1967). Distinguishing theoretical synaptic potentials computed for different soma-dendritic distributions of synaptic input. *Journal of Neurophysiology*, 30(5):1138–1168.

Schneider-Mizell, C. M., Bodor, A. L., Collman, F., Brittain, D., Bleckert, A., Dorkenwald, S., Turner, N. L., Macrina, T., Lee, K., Lu, R., et al. (2023). Cell-type-specific inhibitory circuitry from a connectomic census of mouse visual cortex. *bioRxiv*.

Turner, N. L., Macrina, T., Bae, J. A., Yang, R., Wilson, A. M., Schneider-Mizell, C., Lee, K., Lu, R., Wu, J., Dorkenwald, S., et al. (2022). Reconstruction of neocortex: Organelles, compartments, cells, circuits, and activity. *Cell*, 185(6):1082–1100.

Voruganti, P. (2026). Electrotonic structure shapes partner-level synaptic organization on cortical dendrites. *In preparation*.